

# Attacking and Defending K8s Clusters

NISHAANTH GUNA, MDSEC



## >\$ whoami

- Currently working as a Security Consultant with MDSec, UK focusing on cloud/infrastructure/application security
- 7+ years of penetration testing and security consulting experience
- Previous work experience include Ernst & Young, Mumbai and a mobile security startup in Bangalore conducting mobile pentests
- Reported vulnerabilities to Apple, AT&T, Microsoft, Govt. of UK and many more vulnerabilities in enterprise & popular web/mobile applications
- Chess, photography, hikes
- Conference talks include Steelcon, PHDays, InCTF
- BSides Tirana is my first workshop!

# Agenda

A teal circle with a black outline and a slight drop shadow, containing the text "0x00" in white.

0x00

Containers and  
K8s internals

A teal circle with a black outline and a slight drop shadow, containing the text "0x01" in white.

0x01

Initial access and  
Attack paths

A teal circle with a black outline and a slight drop shadow, containing the text "0x02" in white.

0x02

War stories –  
Common  
misconfigurations

A teal circle with a black outline and a slight drop shadow, containing the text "0x03" in white.

0x03

Securing  
production-  
grade clusters

A teal circle with a black outline and a slight drop shadow, containing the text "0x04" in white.

0x04

VR, CVEs and  
future research

# Docker 101

- Docker automates deployment of software applications inside containers by providing additional layer of abstraction and automation using OS virtualization
- Docker containers != VM
- Containers run as a process in the host operating system segregated by a namespace. Each container gets a different ns.
- Containers provide similar level of isolation at a fraction of computing power.
- Containers run based on the image which you provide. These images can be hosted on public or private repositories (docker.io, ECR and so on)
- Companies managing thousands of docker containers need to orchestrate, maintain all of them at once which is where K8s come in picture.

# Demo

## Objective

- Get familiar with container platform (Docker)
- Run basic commands such as listing containers, executing into the containers and so on
- Understand why there is a need to orchestrate containers

## Commands

- `docker run nginx -dit`
- `docker ps` & `docker ps -a`
- `docker exec -it <identifier> sh`
- `docker container prune`

# Demo

## Objective


- Learn more about docker sockets and misconfigurations
- Under the hood of executing commands inside a docker container and seeing what it reflects in the host OS
- Remap the namespace to use a low-privileged user

## Commands


- `sudo systemctl status docker && ps aux | grep -l dock`
- `docker exec -> run sleep -> check the host OS`
- `cat /etc/subuid && cat /etc/subgid`
- `sudo systemctl stop docker`
- `sudo dockerd --userns-remap="dubov.daniil:dubov.daniil"`

# Kuberwhat??!!

- k(j)ʊ:bər'netɪs, -'neɪtɪs, -'neɪtɪz, -'netɪz
  - koo-ber-net-ees
  - K8s (8 letters between K and s)
  - κυβερνήτης / kubernētēs: Greek for "steersman, navigator" or "guide"
  - Also the etymological root of cybernetics
- 
- Originally designed by Google (Project Borg)
  - Amongst other tech, this is managed by CNCF
  - Open-sourced in 2014
  - What is it though?



BUT THEY COULD NOT UNDERSTAND ITS ALIEN LANGUAGE





# Demo

## Objective


- Setup environmental variables
- Create a multi-node cluster
- Get information about the cluster

## Commands

- NAME=tirbsid.local
- kind create cluster --name \$NAME --config kind-cluster-config.yaml
- less /home/\$USER/.kube/config

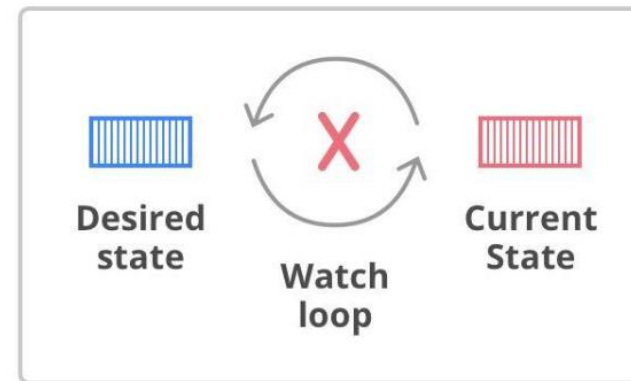
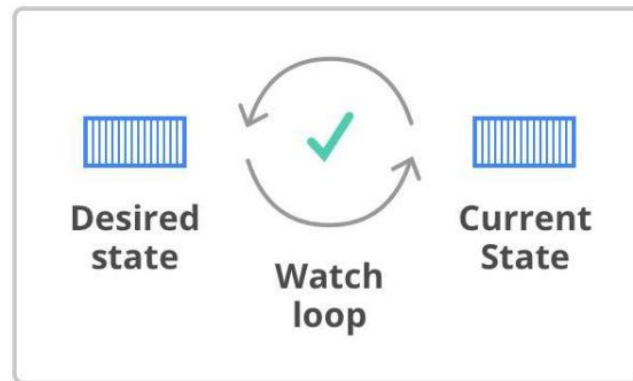
## Why is it sought-after recently?

- Move away from traditional monolith architecture
- K8s uses microservice architecture which results in reduced (almost zero) downtime if used efficiently
- Application Server crashes all of a sudden? K8s got you!
- Easy to orchestrate a large number of containers
- Relies on desired state principle



# Desired state principle

- Any created resource/object will run the exact number of times specified, at any given point of time
- If a container crashes, kube-controller-manager detects this immediately
- When a crash happens and a container goes down, there would be a mismatch between desired state and actual state declared
- A new container will be launched to obtain desired state



# Demo

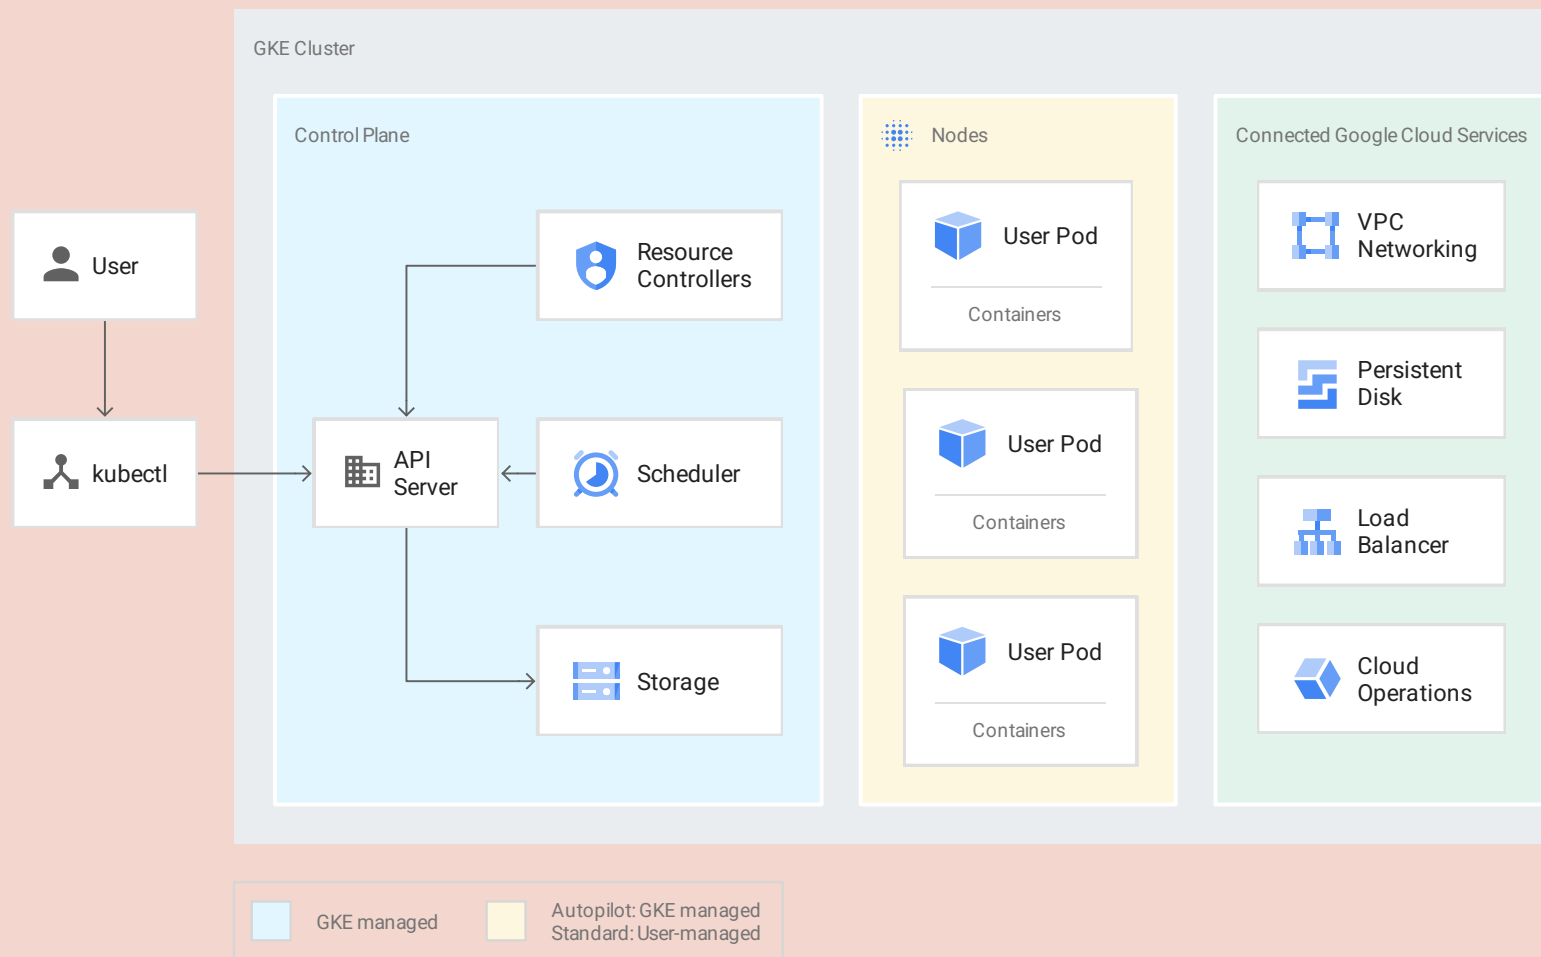
## Objective

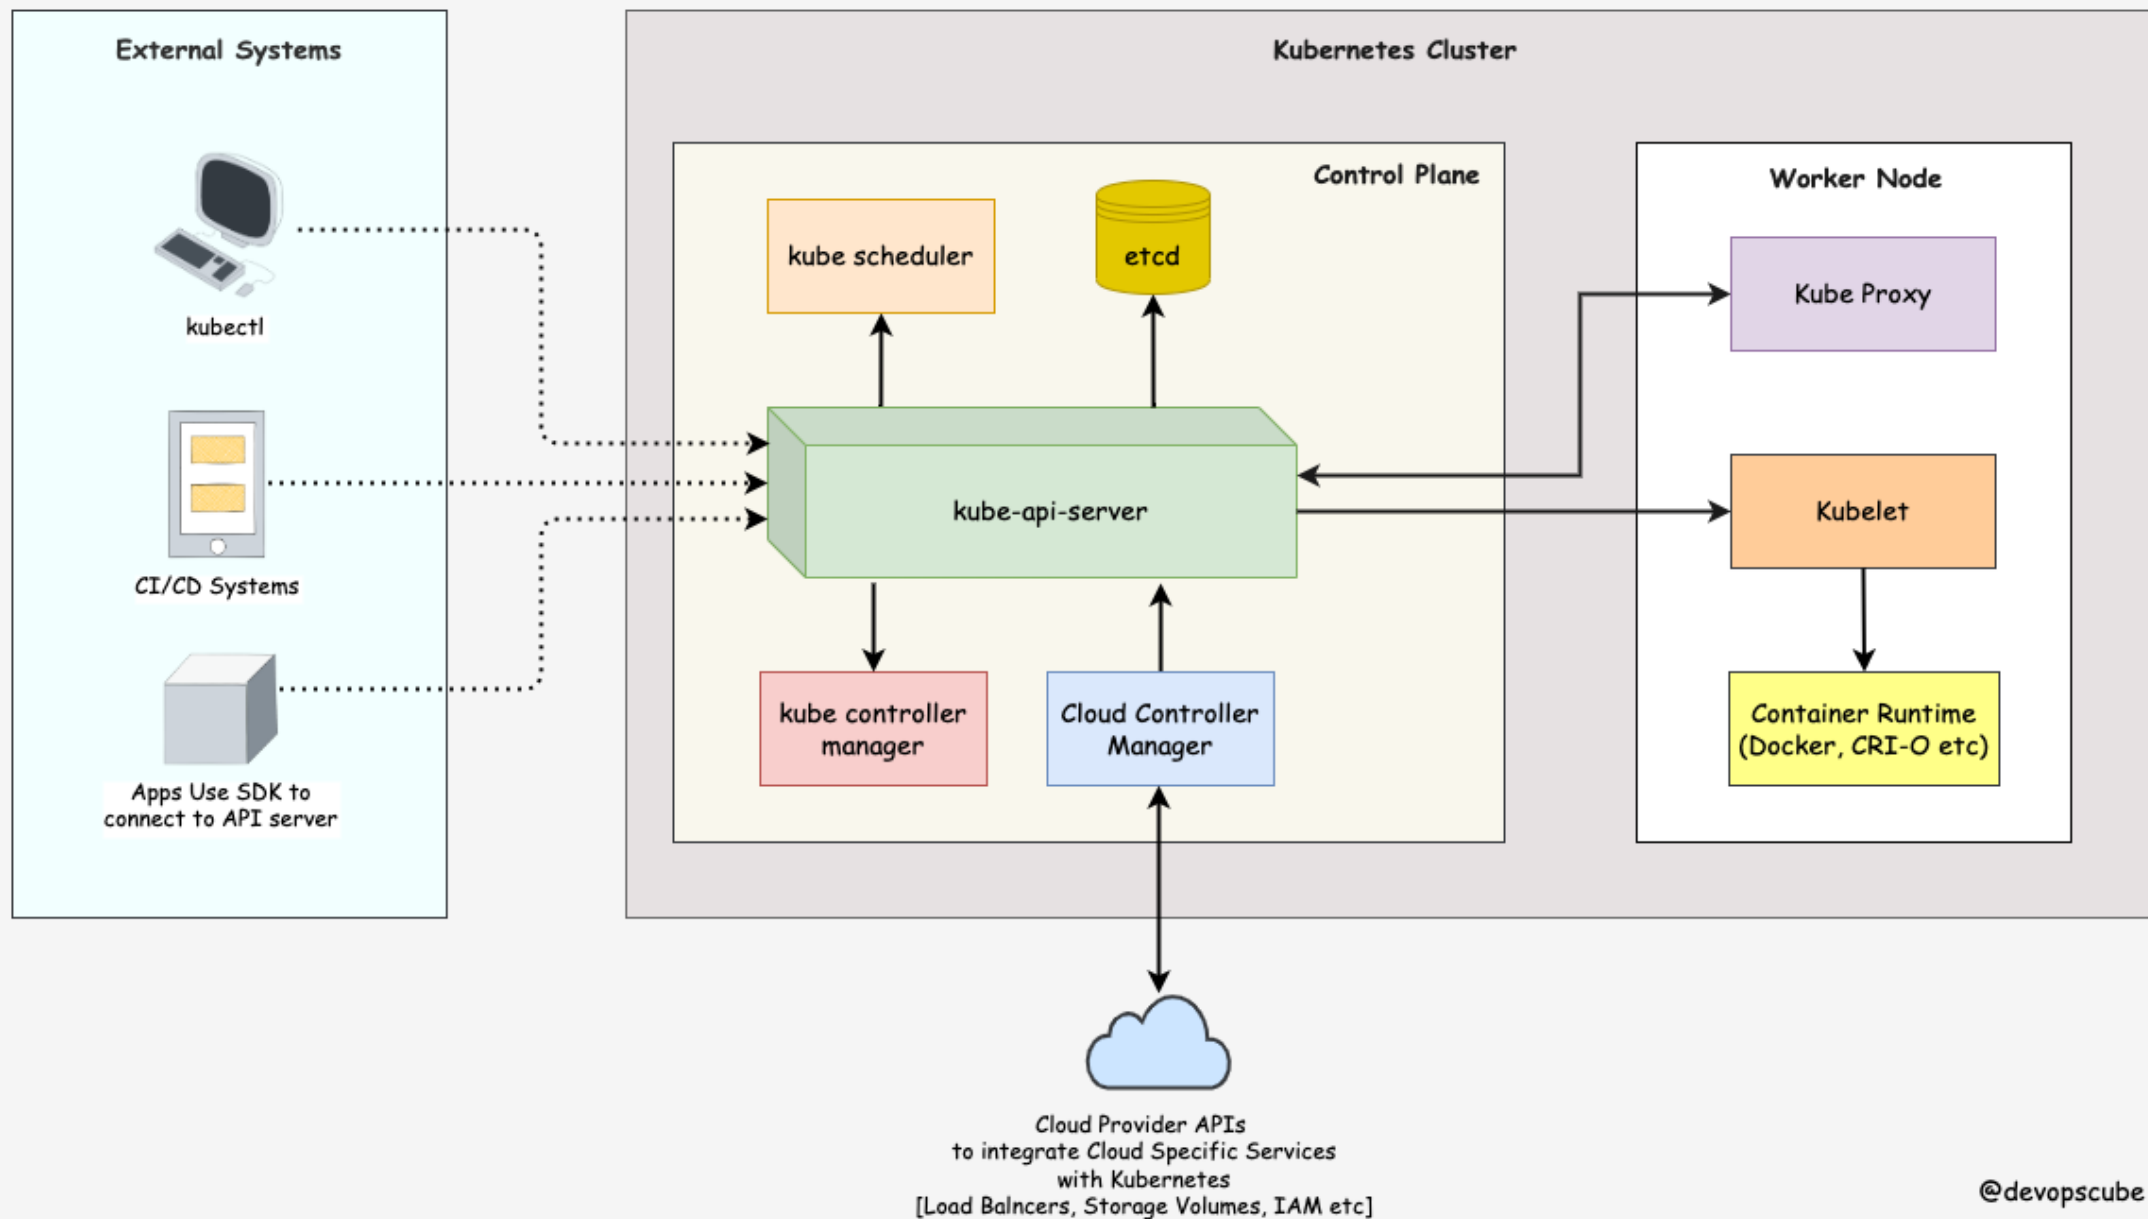
- Understand desired state principle
- Run nginx container as a ReplicaSet
- Crash one of the pods intentionally and see it in effect

## Commands

- `kubectl apply -f nginx-replica.yaml`
- `kubectl delete pod name`
- `kubectl get events | less`





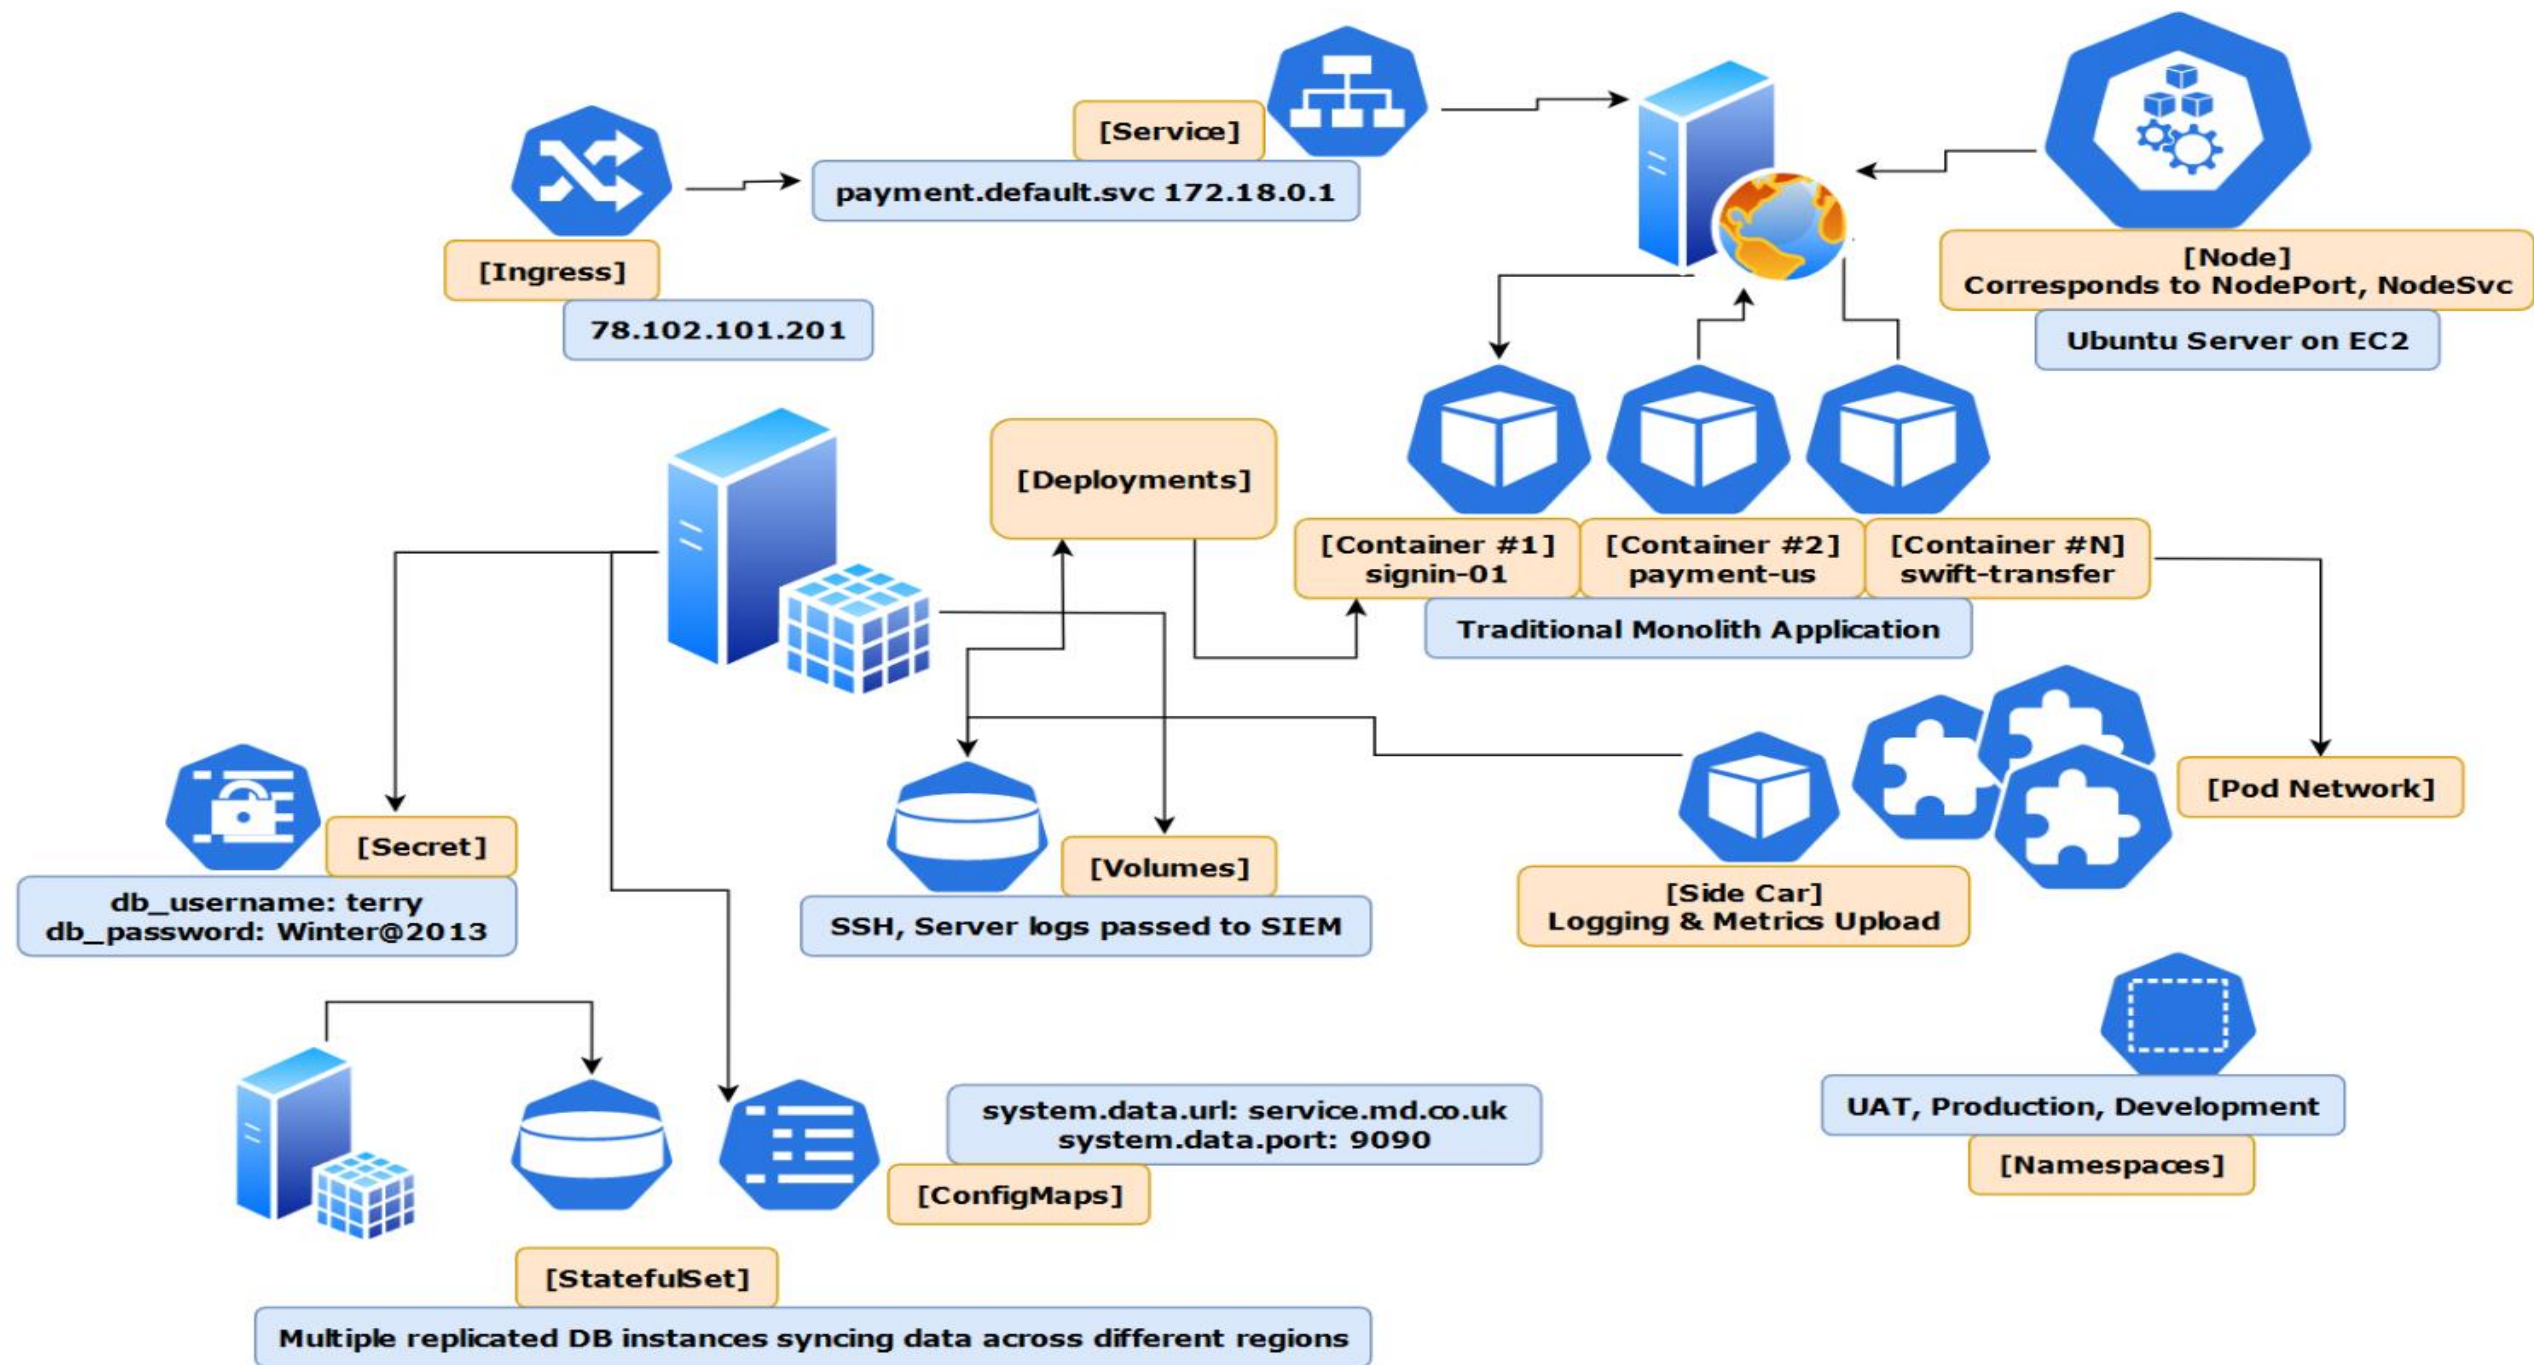


```
$ kubectl api-resources --sort-by name -o wide
```

NAME	SHORTNAMES	APIVERSION	
NAMESPACED	KIND	VERBS	
CATEGORIES			
apiservices		apiregistration.k8s.io/v1	false
APIService		create,delete,deletecollection,get,list,patch,update,watch	
api-extensions			
assign		mutations.gatekeeper.sh/v1	false
Assign		delete,deletecollection,get,list,patch,create,update,watch	
assignimage		mutations.gatekeeper.sh/v1alpha1	false
AssignImage		delete,deletecollection,get,list,patch,create,update,watch	
assignmetadata		mutations.gatekeeper.sh/v1	false
AssignMetadata		delete,deletecollection,get,list,patch,create,update,watch	
bindings		v1	true
Binding		create	
certificatesigningrequests	csr	certificates.k8s.io/v1	false
CertificateSigningRequest		create,delete,deletecollection,get,list,patch,update,watch	
clusterrolebindings		rbac.authorization.k8s.io/v1	false

```
.....
```





# Demo

## Objective

- Create a new namespace to be used for the workshop
- Set aliases for context switching to cluster-admin

## Commands

- `kubectl create ns Tirana`
- `alias switch-admin="kubectl config use-context cluster-admin"`

# Authentication into the Cluster

- Users are of 2 types in K8s – SAs which are managed by K8s and users
- Normal users cannot be added to the cluster through an API call
- Any user that presents a valid certificate signed by the cluster's CA is considered authenticated
- From here, RBAC sub-system kicks in and validates if the authenticated user has privileges on the resource and the namespace
- SAs are users managed by the Kubernetes API, created automatically by API server and bound to a namespace.
- SAs are tied to set of credentials stored as *Secrets*, which are then mounted into pods allowing in-cluster processes to talk to the API Server
- API requests are tied to normal user or a service account or anonymous
- K8s uses client certificates, bearer tokens, or an authenticating proxy to authenticate API requests through plugins

# Authentication into the Cluster

- Auth strategies include X509 certs, static token file, bearer token, bootstrap tokens, SA tokens
- External identity provider can also be used with OpenID connect tokens. This can be integrated with Kerberos, AAD, Salesforce and Google.
- OIDC authentication is rigid way of authenticating into production clusters since the user has to be in the AD, cloud providers (and possible with OTP)
- Webhook token authentication, authenticating proxy can also be used
- Anonymous requests which are part of *system:anonymous* or *system:unauthenticated*
- Flag `--anonymous-auth` has set to be set true in the API server
- User impersonation is possible, impersonate as a user, group or UID

# Demo

## Objective

- Understand the authentication process involved between K8s API Server and kubectl
- Differentiate between SA and certificate authentication

## Commands

- `kubectl get nodes`
- `kubectl get nodes -v9`
- `kubectl get nodes --token $SA`

# RBAC Authorization

- Role-Based Authorization (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users.
- RBAC uses *rbac.authorization.k8s.io* API group
- Role and ClusterRole
- RoleBinding and ClusterRoleBinding
- Decide on the namespace of the resource -> Create a Role or ClusterRole  
-> Bind it a subject -> Validate if RBAC is applied
- Some of the default ClusterRoles include `system:basic-user`, `system:discovery`, `system:public-info-viewer`, `cluster-admin`, `admin`, `edit`, `view`
- Some ClusterRoles are not `system:` prefixed. These are user-facing roles. For example, `cluster-admin`

## Roles and ClusterRole

- An RBAC *Role* or *ClusterRole* contains rules that represent a set of permissions. Permissions are additive (there are no “deny” rules)
- A *Role* sets permissions within the mentioned namespace; whenever you create a role, the namespace has to be mentioned
- *ClusterRole* is a non-namespaced resource. Set of permissions are applied across the cluster
- If you want to define a role within a namespace, use a *Role*; if you want to define a role cluster-wide, use a *ClusterRole*
- *ClusterRole* can be used to grant permissions on cluster-scoped resources like nodes, healthz endpoints)
- Should follow path segment names rule. In other words, the name cannot contain “.” or “..” or “/” or “%”

## RoleBinding and ClusterRoleBinding

- A *RoleBinding* grants permissions defined in a *Role*. It has list of subjects (users, groups or service accounts), and a reference to the role being granted
- A *RoleBinding* grants permissions within a specific namespace whereas a *ClusterRoleBinding* grants it cluster-wide
- A *RoleBinding* may reference any *Role* in the same namespace
- Alternatively, a *RoleBinding* can reference a *ClusterRole* and bind that *ClusterRole* to the namespace of the *RoleBinding*
- If you want to bind a *ClusterRole* to all the namespaces in your cluster, you can use a *ClusterRoleBinding*
- This kind of reference lets you define a set of common roles across the cluster, then reuse them in other namespaces



# Demo

## Objective

- Create a new user in the Kubernetes environment
- Create a new role to allow access to a specific namespace
- Create a new rolebinding to bind the created role to the user

## Commands

- `openssl genrsa -out dubov.key 2048`
- `openssl req -new -key dubov.key -out dubov.csr -subj "/CN=dubov"`
- `docker cp container-name:/etc/kubernetes/pki/ca.key .`
- `docker cp fc25868e5b43:/etc/kubernetes/pki/ca.crt .`

# Demo

## Commands

- openssl x509 -req -in dubov.csr -CA ca.crt -CAkey ca.key -CAcreateserial -out dubov.crt -days 300
- kubectl config set-credentials dubov --client-certificate=dubov.crt --client-key=dubov.key
- kubectl config set-context dubov --cluster=kind-bsides --user=dubov --server=https://127.0.0.1:35605

## Commands

- kubectl config use-context dubov
- kubectl get pods & kubectl auth can-i --list
- kubectl apply -f rbac-roles.yaml
- kubectl apply -f rbac-rolebinding.yaml

# Demo

## Commands

- Dump necessary files for roles, cluster role and corresponding bindings
- Identify risky RBAC permissions using rbac-audit
- Identify attack paths (plugins, resources) using rbac-police

## Commands

- `kubectl get roles -A -o json > roles.json`
- `kubectl get clusterroles -A -o json > clusterroles.json`
- `kubectl get rolebindings -A -o json > rolebindings.json`
- `kubectl get clusterrolebindings -A -o json > clusterrolebindings.json`

```
➤ python3 ExtensiveRoleCheck.py --clusterRole /tmp/clusterroles.json --role /tmp/roles.json --clusterrolebindings /tmp/clusterrolebindings.json --rolebindings /tmp/rolebindings.json
```

```
[*] Started enumerating risky ClusterRoles:  
[!][ClusterRole]→ gatekeeper-manager-role Has permission to use list on any resource!  
[!][ClusterRole]→ gatekeeper-manager-role Has permission to use delete on any resource!  
[!][ClusterRole]→ gatekeeper-manager-role Has permission to use delete on any resource!  
[!][ClusterRole]→ gatekeeper-manager-role Has permission to use delete on any resource!  
[!][ClusterRole]→ gatekeeper-manager-role Has permission to use delete on any resource!  
[!][ClusterRole]→ get-secrets-everywhere Has permission to list secrets!  
[!][ClusterRole]→ local-path-provisioner-role Has permission to access pods with any verb!  
[*] Started enumerating risky Roles:  
[!][Role]→ create-pod-ns Has permission to create pods!  
[!][Role]→ gatekeeper-manager-role Has permission to list secrets!  
[!][Role]→ create-pod-ns Has permission to create pods!  
[*] Started enumerating risky ClusterRoleBinding:  
[!][ClusterRoleBinding]→ gatekeeper-manager-rolebinding is binded to gatekeeper-admin ServiceAccount.  
[!][ClusterRoleBinding]→ local-path-provisioner-bind is binded to local-path-provisioner-service-account Service Account.  
[!][ClusterRoleBinding]→ rbac-clusterrole-dubov is binded to the User: dubov!  
[*] Started enumerating risky RoleBindings:  
[!][RoleBinding]→ rbac-role-binding-role-binding is binded to the User: dubov!  
[!][RoleBinding]→ rbac-role-dubov is binded to the User: dubov!  
[!][RoleBinding]→ gatekeeper-manager-rolebinding is binded to gatekeeper-admin ServiceAccount.  
➤ ggwp@metaaoh:~/k8s-content/kubernetes-rbac-audit [master]➤
```

➤ pts/0 ➤

—<ggwp@metaaoh:~/k8s-content/rbac-police [main]>—

—<pts/0>—

—<%>— rbac-police eval lib/

```
{
  "policyResults": [
    {
      "policy": "lib/assign_sa.rego",
      "severity": "Critical",
      "description": "Identities that can create pods or create, update or patch pod controllers (e.g. DaemonSets, Deployments, Jobs) in privileged namespaces (kube-system), may assign an admin-equivalent SA to a pod in their control",
      "violations": {
        "serviceAccounts": [
          {
            "name": "local-path-provisioner-service-account",
            "namespace": "local-path-storage",
            "nodes": [
              {
                "bsides-control-plane": [
                  "local-path-provisioner-6bc4bddd6b-xr4ks"
                ]
              }
            ]
          }
        ]
      }
    }
  ]
}
```



```

{
  "policy": "lib/list_secrets.rego",
  "severity": "Medium",
  "description": "Identities that can list secrets cluster-wide may access confidential information, and in some cases serviceAccount tokens",
  "violations": {
    "serviceAccounts": [
      {
        "name": "gatekeeper-admin",
        "namespace": "gatekeeper-system",
        "nodes": [
          {
            "bsides-control-plane": [
              "gatekeeper-audit-6668847c5c-549wz",
              "gatekeeper-controller-manager-7fff77f764-c2bm7",
              "gatekeeper-controller-manager-7fff77f764-nkmml",
              "gatekeeper-controller-manager-7fff77f764-qv2sl"
            ]
          }
        ]
      }
    ]
  }
},

```

# K8s Vulns Overview



0x00

Deploying  
resources in  
default  
namespace



0x01

Bad Pods –  
Container to Host  
Compromise



0x02

Malicious  
Container  
Images  
(Vulnerable  
CI/CD, Jenkins)



0x03

Insecure RBAC  
(\*.\* on resources,  
verbs)



0x04

Overly-permissive  
and powerful  
third-party plugins

## Initial Access

- Exposed cloud credentials – Compromised cloud credentials used by exposed AKS/GKE/EKS can lead to cluster takeover if keys are scoped
- Compromised image in registry – Compromising a private registry to plant malware in the base image pulled regularly by developers
- Exposed Kubeconfig – If attackers get access to kubeconfig file via a compromised client or any others means, this can be used to access K8s.
- Application vulnerability – Running a public-facing vulnerable application in a cluster which is vulnerable to RCE, LFI can lead to compromise. An attacker can read the SA mounted into the pod, breakout to attack the host
- Exposed Interfaces – Software like K8s dashboard, Apache NiFi, Kubeflow, Argo were never meant to be exposed publicly.



# Demo

## Objective

- Understand why deploying resources in default namespace is insecure
- Understand how SAs integrate into the K8s environment
- Exploit privileged SA to list secrets on the cluster

## Commands

- `kubectl run nginx --image=nginx`
- `kubectl exec -it nginx -- bash`
- `cat /var/run/secrets/kubernetes.io/serviceaccount/token`
- `kubectl auth can-i --list --token $SA`
- `kubectl get nodes`

# Demo

## Objective

- Understand why running pods with host mounts can be dangerous in an environment

## Commands

- `docker exec -it id bash`
- `kubectl run nginx --image=nginx`
- `kubectl apply -f kryptopod.yaml`
- `kubectl exec -it hot-pod -- bash`

# Privilege Escalation – Attack Paths

## Scenario – 1

- {-} Compromised credentials has untethered permissions on the entire cluster
- {-} Get the list of nodes running in the environment to differentiate between master and worker nodes
- {-} Create a pod with all the sensitive host mounts, capabilities and PrivEsc
- {-} Using nodeName value, deploy the pod in one of the master nodes
- {-} From inside the privileged pod, chroot to the host filesystem and search through the files and folders in the system
- {-} Environmental variable contains a Service Account token and location to kubeconfig file
- {-} Steal them to get cluster-admin
- {-} Pwn!

# Privilege Escalation – Attack Paths

## Scenario – 2

- {-} Cordon preventing unauthorized pods in the master nodes
- {-} Similar to the previous environment, create a pod on the worker node instead of the master node
- {-} From the container, chroot inside the worker node and enumerate all the pods/namespaces running using `ps -ef`
- {-} Write a script to enter all the namespaces and save the SA token
- {-} One of the pods running database manager had super privileges due to the nature of the pod. Was one of the centralized key-value managers
- {-} Use the database manager pod's SA token to obtain cluster-admin
- {-} Pwn!

# Privilege Escalation – Attack Paths

## Scenario – 3

- {-} Internal K8s assessment with OIDC login to the cluster
- {-} Production environment completely locked out rendering the stolen keys completely useless.
- {-} Key only has GET, LIST permissions on a few resources
- {-} Take a step back and view the kubeconfig file once again
- {-} Notice multiple contexts in the configuration file. One of the context points to staging cluster
- {-} Re-use the keys in the staging cluster to discover the same keys have extra privileges including CREATE pods on lot of sensitive namespaces
- {-} Create a pod with hostPath flag to steal the cluster-admin credentials
- {-} Pwn!

# Privilege Escalation – Attack Paths

## Scenario – 4

- {-} Initial access through Local File Read
- {-} Compromised pod has SA with the ability to CREATE pods in a specific ns
- {-} Pods cannot run as privileged due to PodSecurityPolicy.  
Pods cannot run be deployed with hostNetwork/hostPath flags due to additional security controls in the environment to prevent breakouts.
- {-} Can be bypassed by creating a pod with hostPID or hostIPC flags
- {-} Using hostPID, list the processes on the host machine
- {-} Identify a process related to cloud environment running with AWS access and secret keys as arguments
- {-} AWS keys had access to the entire cloud estate == Pwn!

# Privilege Escalation – Attack Paths

## Scenario – 5

- {-} Initial Access through SSRF from a web application
- {-} Request the AWS Instance Metadata service (169.254.169.254)
- {-} Keys are privileged due to a design flaw in AWS environment. Using the obtained keys you can pull ECR images and list few S3 buckets
- {-} Export the keys to the local machine and verify the identity of the key
- {-} Enumerate S3 buckets of the organization and use the stolen key to fetch the objects of the bucket
- {-} One of the objects in the bucket include ETCD backup of the prod cluster
- {-} Extract the ETCD store to discover a trove of SA tokens and AWS keys
- {-} One of the keys had cluster-admin access in the cluster
- {-} Pwn!

## AWS EKS PrivEsc w SSRF

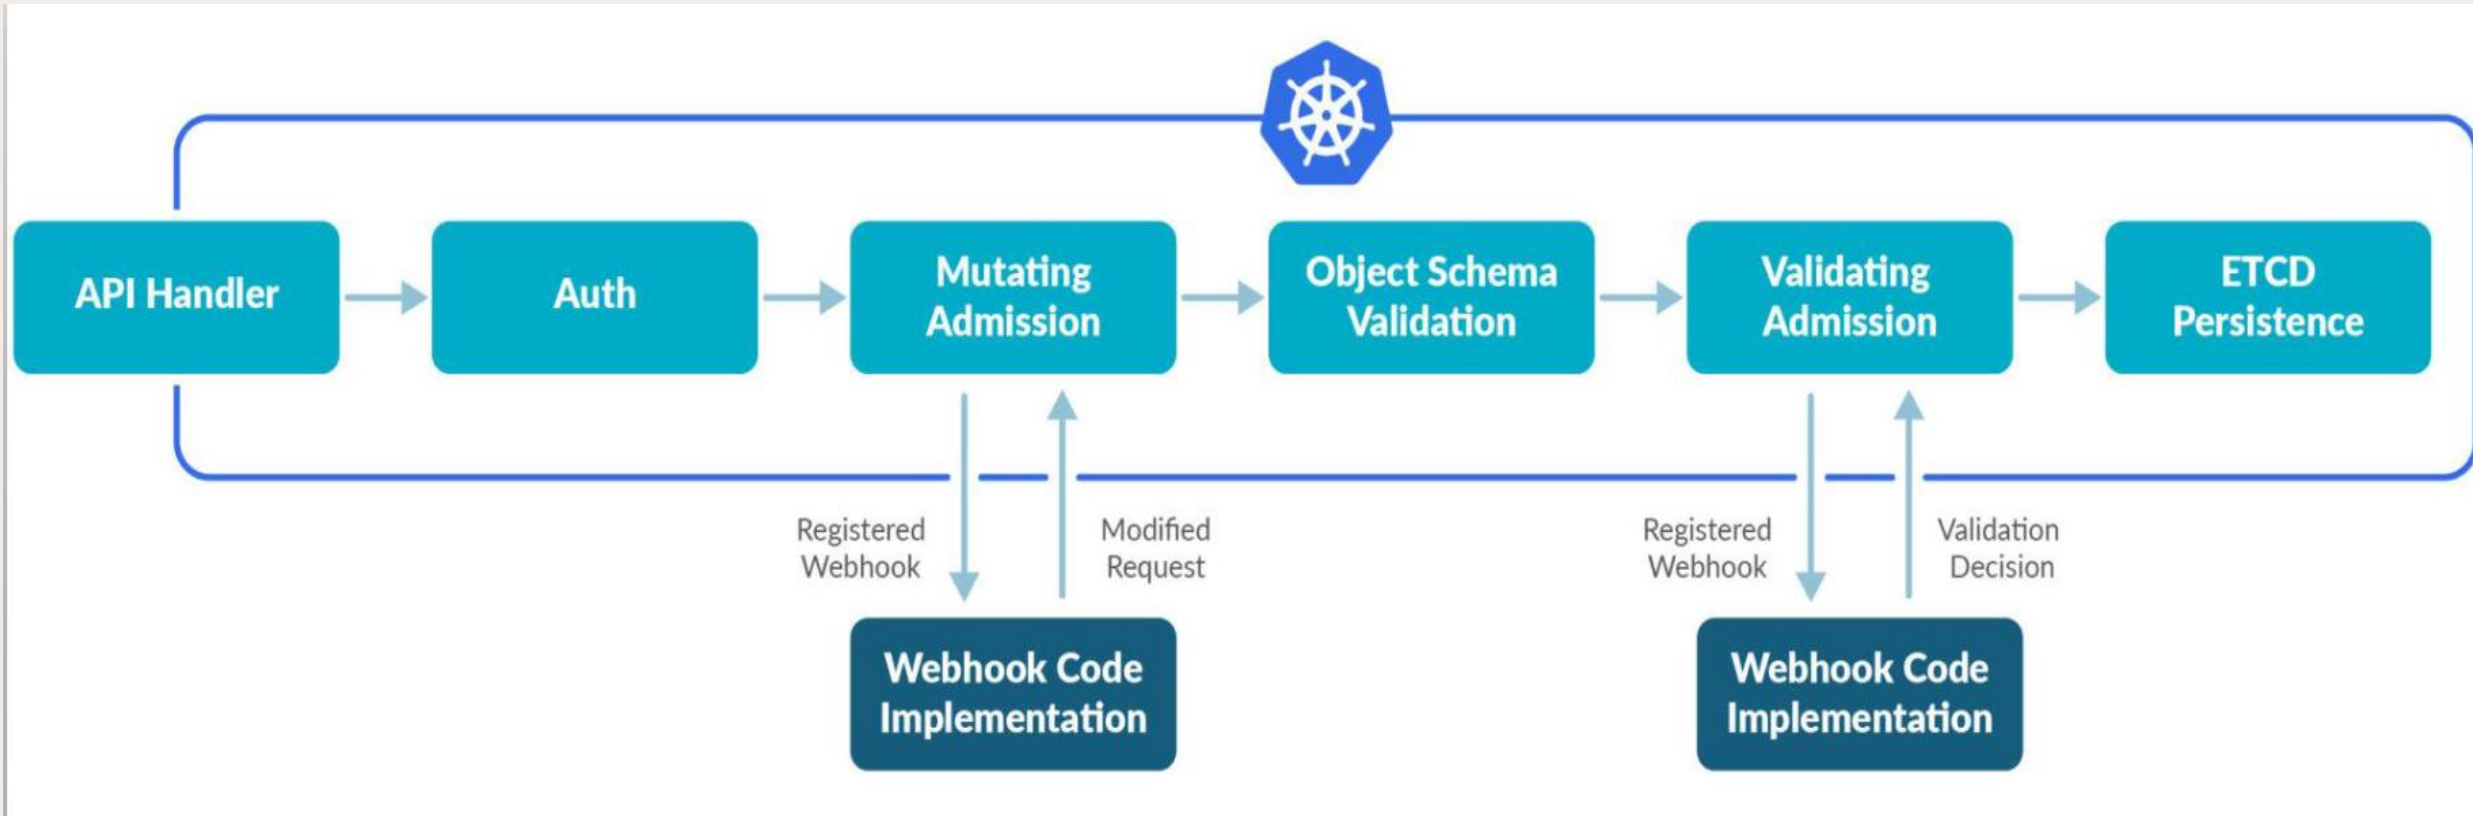
- Able to reach AWS Metadata API from inside a container running on a node and no default mechanism to stop that
- On a low-privileged EKS pod, the temporary AWS keys has same level of privilege as underlying EC2 machine
- Compromising a single pod using the keys from 169.254.169.254 can lead to potential compromise of AWS estate
- Block pods from reaching the AWS Metadata API
- Issue can be fixed by using IMDSv2
- Similar issues on GKE, but detailed documentation on how to prevent this vulnerability



# Admission Controllers

- Gatekeeps the cluster from object writes to ETCD server if a resource does not conform with a policy (e.g., reject images not pulled from internal.mdsec.docker.registry.com, do not run containers as root or with privileged flag)
- Possible to create HTTP callbacks/webhooks with custom logic to decide if a resource should be let in the cluster or not
- Mutating and Validating Controllers
- Mutating Controllers can intercept an API request and PATCH the request to manipulate a struct or an object in it)
- 30+ shipped with K8s, compiled into the kube-apiserver binary (NamespaceLifecycle, PodSecurityPolicy, ValidatingAdmissionWebhook)
- Flask Server monitoring for webhooks -> Deploy it as a SVC -> Register a Webhook Controller -> Validate if resources conform to policy

# Admission Controllers - Flow



```
// AdmissionReview describes an admission review request/response.
▼ type AdmissionReview struct {
    metav1.TypeMeta `json:",inline"`
    // Request describes the attributes for the admission request.
    // +optional
    Request *AdmissionRequest `json:"request,omitempty" protobuf:"bytes,1,opt,name=request"`
    // Response describes the attributes for the admission response.
    // +optional
    Response *AdmissionResponse `json:"response,omitempty" protobuf:"bytes,2,opt,name=response"`
}
```

```

type AdmissionRequest struct {
    UID types.UID `json:"uid" protobuf:"bytes,1,opt,name=uid"`
    Kind metav1.GroupVersionKind `json:"kind" protobuf:"bytes,2,opt,name=kind"`
    Resource metav1.GroupVersionResource `json:"resource" protobuf:"bytes,3,opt,name=resource"`
    SubResource string `json:"subResource,omitempty" protobuf:"bytes,4,opt,name=subResource"`
    RequestKind *metav1.GroupVersionKind `json:"requestKind,omitempty"
protobuf:"bytes,13,opt,name=requestKind"`

    RequestResource *metav1.GroupVersionResource `json:"requestResource,omitempty"
protobuf:"bytes,14,opt,name=requestResource"`
    RequestSubResource string `json:"requestSubResource,omitempty"
protobuf:"bytes,15,opt,name=requestSubResource"`

    Name string `json:"name,omitempty" protobuf:"bytes,5,opt,name=name"`
    Namespace string `json:"namespace,omitempty" protobuf:"bytes,6,opt,name=namespace"`
    Operation Operation `json:"operation" protobuf:"bytes,7,opt,name=operation"`

    UserInfo authenticationv1.UserInfo `json:"userInfo" protobuf:"bytes,8,opt,name=userInfo"`

    Object runtime.RawExtension `json:"object,omitempty" protobuf:"bytes,9,opt,name=object"`

    OldObject runtime.RawExtension `json:"oldObject,omitempty" protobuf:"bytes,10,opt,name=oldObject"`

    DryRun *bool `json:"dryRun,omitempty" protobuf:"varint,11,opt,name=dryRun"`
    Options runtime.RawExtension `json:"options,omitempty" protobuf:"bytes,12,opt,name=options"`
}

```

```

type AdmissionResponse struct {
    // UID is an identifier for the individual request/response.
    // This should be copied over from the corresponding AdmissionRequest.
    UID types.UID `json:"uid" protobuf:"bytes,1,opt,name=uid"`

    Allowed bool `json:"allowed" protobuf:"varint,2,opt,name=allowed"`

    Result *metav1.Status `json:"status,omitempty" protobuf:"bytes,3,opt,name=status"`

    Patch []byte `json:"patch,omitempty" protobuf:"bytes,4,opt,name=patch"`

    // The type of Patch. Currently we only allow "JSONPatch".
    // +optional
    PatchType *PatchType `json:"patchType,omitempty" protobuf:"bytes,5,opt,name=patchType"`

    // AuditAnnotations is an unstructured key value map set by remote admission controller (e.g. error=image-
    blacklisted).
    // MutatingAdmissionWebhook and ValidatingAdmissionWebhook admission controller will prefix the keys with
    // admission webhook name (e.g. imagepolicy.example.com/error=image-blacklisted). AuditAnnotations will be
    provided by

    AuditAnnotations map[string]string `json:"auditAnnotations,omitempty"
    protobuf:"bytes,6,opt,name=auditAnnotations"`

    Warnings []string `json:"warnings,omitempty" protobuf:"bytes,7,rep,name=warnings"`
}

```

```
apiVersion: admissionregistration.k8s.io/v1
kind: ValidatingWebhookConfiguration
metadata:
  name: "pod-policy.example.com"
webhooks:
- name: "pod-policy.example.com"
rules:
- apiGroups:  [""]
  apiVersions: ["v1"]
  operations: ["CREATE"]
  resources:  ["pods"]
  scope:      "Namespaced"
  clientConfig:
    service:
      namespace: "example-namespace"
      name: "example-service"
  caBundle: <CA_BUNDLE>
  admissionReviewVersions: ["v1"]
  sideEffects: None
```

# Demo

## Objective

- Understand how Admission Controllers work
- Check what plugins are enabled in our cluster by default
- Deploy Kyverno/OPA Gatekeeper and check validation

## Commands

- `kubectl describe pod kube-apiserver-bsides-control-plane -n kube-system`
- `ls -la /etc/kubernetes/manifests`
- `kubectl apply -f https://raw.githubusercontent.com/open-policy-agent/gatekeeper/master/deploy/gatekeeper.yaml`
- `kubectl get ns`

# Demo

## Commands

- `kubectl apply -f https://raw.githubusercontent.com/open-policy-agent/gatekeeper-library/master/library/pod-security-policy/privileged-containers/template.yaml`
- `kubectl apply -f priv-constraint.yaml`
- `kubectl run nginx-opa-test --image nginx --privileged=true`

## Commands

- `kubectl apply -f https://raw.githubusercontent.com/open-policy-agent/gatekeeper-library/master/library/general/allowedrepos/template.yaml`
- `kubectl apply -f imagepull-constraint.yaml`
- `kubectl run image-opa-test --image nginx`



## ETCD Best Practices

- Due to the nature of ETCD, it is a critical resource. Even a read access to the daemon is dangerous because it contains sensitive information about the secrets, state of the cluster.
- Enabling encryption at rest for ETCD secrets. Switched off by default. Enable it by using `--encryption-provider-config`
- Isolate the daemon behind a firewall or use network ACL to limit communication between pods and the key-value store
- Backups should be done regularly and stored in a secure location
- Audit logging and monitoring for malicious events, unauthorized entries
- Update all the components in the environment regularly and keep a look out for CVEs which affect the software in your environment.

# Setup Multi-Node Cluster w Vagrant

```
Vagrantfile
14  Vagrant.configure("2") do |config|
15
16      config.vm.define "k8s-master" do |master|
17          master.vm.box = IMAGE_NAME
18          master.vm.network "private_network", ip: API_SERVER_IP
19          master.vm.hostname = "#{CLUSTER_NAME}-k8s-control-plane"
20          master.vm.provider :virtualbox do |vb|
21              vb.name = "#{CLUSTER_NAME}-k8s-control-plane"
22              vb.memory = VM_MEMORY
23              vb.cpus = 2
24          end
25          master.vm.provision "master-common", type: "shell",
26              env: {
27                  "API_SERVER_IP": API_SERVER_IP,
28                  "DNS_SERVER": DNS_SERVER,
29                  "KUBERNETES_VERSION": KUBERNETES_VERSION,
30                  "CRI_VERSION": CRI_VERSION,
31                  "CLUSTER_NAME": CLUSTER_NAME,
32                  "OS": OS
33              },
34              path: "init/common.sh"
```

# Setup Multi-Node Cluster w Vagrant

```
init > $ common.sh
44 apt update -y
45 apt install cri-o cri-o-runc -y
46 systemctl daemon-reload
47 systemctl enable crio --now
48
49 if [ ! -d /etc/modules.load.d/ ]; then
50 |   mkdir /etc/modules.load.d/
51 | fi
52 cat <<EOF | sudo tee /etc/modules.load.d/crio.conf
53 overlay
54 br_netfilter
55 EOF
56
57 modprobe overlay
58 modprobe br_netfilter
59
60 cat <<EOF | sudo tee /etc/sysctl.d/99-kubernetes-cri.conf
61 net.bridge.bridge-nf-call-iptables = 1
62 net.ipv4.ip_forward                 = 1
```

# Setup Multi-Node Cluster w Vagrant

```
init > $ master.sh
1  #!/bin/bash
2  #
3  set -euxo pipefail
4
5  # control plane init and image pull
6  kubeadm config images pull
7  kubeadm init --apiserver-advertise-address=$API_SERVER_IP --apiserver-cert-extra-sans=
8  echo "[-+--+--+--+--+!K8S Control Panel Up!-+--+--+--+--+]"
9
10 # copy token script and kubeconfig from master
11 kubeadm token create --print-join-command | tee /vagrant/node_join.sh
12 cp /etc/kubernetes/admin.conf /vagrant/admin.conf
13 echo "[-+--+--+--+--+!K8S Credentials copied to host OS /vagrant folder!-+--+--+--+--+]"
14
```

# Setup Multi-Node Cluster w Vagrant

```
0e478
k8s-master: + echo '[-+--+--+--+--+!K8S Control Panel Up!-+--+--+--+--+]'
k8s-master: [-+--+--+--+--+!K8S Control Panel Up!-+--+--+--+--+]
k8s-master: + tee /vagrant/node_join.sh
k8s-master: + kubeadm token create --print-join-command
k8s-master: kubeadm join 192.168.13.37:6443 --token xel22w.gn1ccd2d8n9q5oxy --discovery-
810ef29a5d70491114fa7633636cfb0862d0cf50abfa97a9cde253b40e478
k8s-master: + cp /etc/kubernetes/admin.conf /vagrant/admin.config
k8s-master: [-+--+--+--+--+!K8S Credentials copied to host OS /vagrant folder!-+--+--+--+--+]
k8s-master: + echo '[-+--+--+--+--+!K8S Credentials copied to host OS /vagrant folder!-+--+--+--+--+]'
k8s-master: + mkdir -p /home/vagrant/.kube
k8s-master: + cp /etc/kubernetes/admin.conf /home/vagrant/.kube/config
k8s-master: + chown -R vagrant: /home/vagrant
k8s-master: [-+--+--+--+--+!K8S Credentials exported to /home/vagrant/.kube/!-+--+--+--+--+]
k8s-master: + echo '[-+--+--+--+--+!K8S Credentials exported to /home/vagrant/.kube/!-+--+--+--+--+]'
==> k8s-worker-1: Importing base box 'bento/ubuntu-22.04'...
==> k8s-worker-1: Matching MAC address for NAT networking...
==> k8s-worker-1: Checking if box 'bento/ubuntu-22.04' version '202303.13.0' is up to date..
==> k8s-worker-1: Setting the name of the VM: enron.corp.k8s.local-k8s-worker-node-1
==> k8s-worker-1: Fixed port collision for 22 => 2222. Now on port 2200.
==> k8s-worker-1: Clearing any previously set network interfaces...
```

## Try it out!

[ ] <https://github.com/nishaanthguna22/vagrant-deployment-files/tree/master/K8s-multicluster>

[ ] <https://t.ly/bVMv>



# In a Nutshell

- RBAC is critical for resource authz in the cluster
- Container breakouts often result in host compromise.
- Third-party plugins hoard critical vulnerabilities due to the excessive privileges on the SAs
- Compromise the pod with super-privileged SA equates to cluster-admin
- Do not expose the cluster's API Server to the Internet without any firewall or network restrictions
- Certificate authentication is best suited for enterprise-grade clusters



# References

- <https://imgflip.com/>
  - <https://twitter.com/memenetes>
  - <https://cloud.google.com>
  - <https://www.adaltas.com/en/2019/08/07/users-rbac-kubernetes/>
  - <https://kubernetes.io>
  - <https://microsoft.github.io/Threat-Matrix-for-Kubernetes>
  - <https://blog.christophetd.fr/privilege-escalation-in-aws-elastic-kubernetes-service-eks-by-compromising-the-instance-role-of-worker-nodes/>
  - <https://blog.calif.io/p/privilege-escalation-in-eks>
  - <https://faun.pub/top-20-kubernetes-memes-b5cb4c5af395>
  - <https://bishopfox.com/blog/kubernetes-pod-privilege-escalation>
-